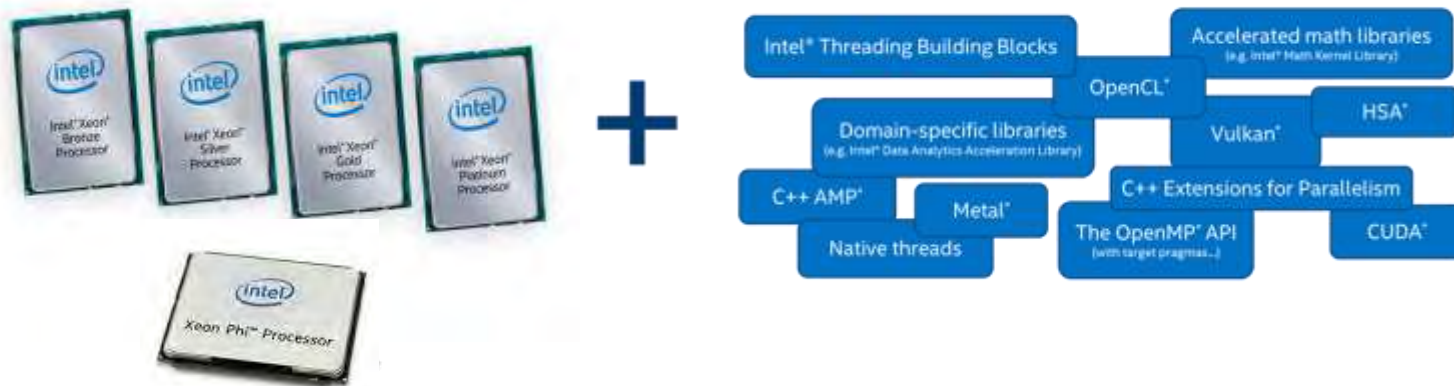




# SOLVING PERFORMANCE PUZZLE BY USING INTEL® THREADING BUILDING BLOCKS

Evgeny Fiksmen  
Intel Corporation  
October 2017

# CHALLENGES IN PROGRAMING OF MODERN SYSTEMS




## ■ Hardware Progress in 6 years

- Over 3x more cores, 8 (Intel® Xeon E5) vs 28 (Intel® Xeon Platinum 8180) cores, or 72 (Intel® Xeon Phi™ 2<sup>nd</sup> gen)
- 4x wider registers , Intel® SSE4.2 vs Intel® AVX-512
- Under same power envelope -> Lower frequency
- Heterogeneity

# THE THREE LEVELS OF PARALLELISM

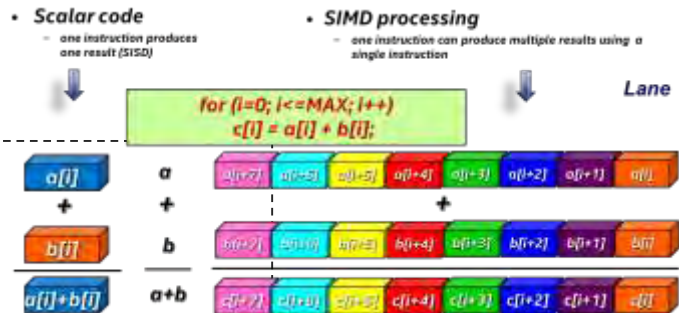
- Task / Message Driven
  - GRID engines, MPI, Intel TBB Flow Graph
- Fork-Join Parallel Regions
  - OpenMP\*
  - Intel TBB
  - Parallel STL in C++17(Requires runtime support)
- (auto-)Vectorization / SIMD
  - Libraries
  - Intrinsics
  - OpenMP 4.x(Requires compiler support)
  - Parallel STL in C++17(Requires compiler support)



# VECTORIZATION

```
void foo(const double* a, const double* b, double* c, int n)
{
    for(int i=0; i<n; ++i)
        c[i] = bar(a[i],b[i]);
}
```

“Serial brick”



```
void foo(const double* a, const double* b, double* c, int n)
{
    #pragma omp simd
    for(int i=0; i<n; ++i)
        c[i] = bar(a[i],b[i]);
}
```

“vectorization”

```
#pragma omp declare simd
double bar(double a, double b)
{ return sqrt(a/b); }
```

Vector function


```
...
void foo(const double* a, const double* b, double* c, int n)
{
    #pragma omp simd
    for(int i=0; i<n; ++i)
        c[i] = bar(a[i], b[i]);
}
```

# MEMORY LAYOUT OPTIMIZATION – INTEL® SDLT


## SIMD DATA LAYOUT TEMPLATES

- A C++11 template library providing concepts of Containers, Accessors and Offsets
  - **Containers** encapsulate the in memory data layout of an Array of “Plain Old Data” objects.
  - SIMD loops use **accessors** with an array subscript operator (just like C++ arrays) to read from or write to the objects in the Containers.
  - **Offsets** can be embedded in accessors or applied to a Index passed to the accessors array subscript operator.
- Usages:
  - Image processing
  - 2D PDE solvers
  - Interpolators

### AOS in Memory Layout



### SOA in Memory Layout



# INTEL SDLT EXAMPLE

```
using namespace sdtl;
struct YourStruct
{
    float x;
    float y;
    float z;
};
SDLT_PRIMITIVE(YourStruct, x, y, z)
...
auto shape = n_extents[128][256][512];
typedef n_container<YourStruct, layout::soa, decltype(shape)> Container3d;

Container3d input(shape), output(shape);
auto inputs = input.const_access();
auto outputs = output.access();

for(int z = 0; z < shape.extent_d0(); ++z) {
    for(int y = 0; y < shape.extent_d1(); ++y) {
        #pragma omp simd
        for(int x = 0; x < shape.extent_d2(); ++x) {
            YourStruct val = inputs[z][y][x];
            YourStruct result = ... // compute result
            outputs[z][y][x] = result;
        }
    }
}
```

C++ lacks compile time reflection, so the user must provide SDLT with some information on the layout of YourStruct.

# FORK-JOIN: INTEL® THREADING BUILDING BLOCKS (INTEL® TBB)




- A widely used C++ template library for parallel programming
- What
  - Task based scheduling
  - Parallel algorithms and data structures
    - `parallel_for`, `parallel_reduce`, ...
    - `concurrent_hash`, `concurrent_queue`, ...
  - Dependency graphs and data flow primitives
  - Thread and synchronization primitives
    - `atomics`, `TLS`, `mutex`, ...
  - Scalable memory allocation – efficient also for non threaded applications
- Benefits
  - Is a library-only solution that does not depend on compiler support
  - Is both a commercial product and an open-source (Apache v2.0)
  - Supports C++, Windows\*, Linux\*, OS X\*, Android\* and other OSes
  - Commercial support for Intel® Atom™, Core™, Intel® Xeon® & Intel® Xeon Phi™ processors

# FORK-JOIN:PARALLEL REGIONS

```
#pragma omp declare simd
double bar(double a, double b)
{ return sqrt(a/b); }
void foo(const double* a, const double* b, double* c, int n)
{
    #pragma omp simd
    for(int i=0; i<n; ++i)
        c[i] = bar(a[i], b[i]);
}
```

“vectorized brick”



```
#pragma omp declare simd
double bar(double a, double b)
{ return sqrt(a/b); }

tbb::parallel_for(tbb::blocked_range<int>(0,n,STEP),
    [&](const tbb::blocked_range<int>& r)
    {
        auto begin = r.begin(), end = r.end();
        #pragma omp simd
        for(auto ii=begin; ii<end; ++ii)
            c[ii] = bar(a[ii], b[ii]);
    }
);
```

C++11: Lambda  
Expression


OpenMP\*4.0 simd directive co-exists with a different multi-threading framework




# TASK PARALLISM: INTEL TBB FLOW GRAPH

- Introduced in Intel TBB 4.0
- Efficient implementation of dependency graph and data flow algorithms
- Initially Designed for shared memory application
- Enables developers to exploit parallelism at higher levels

Hello World




```
graph g;
continue_node< continue_msg > h( g,
    []( const continue_msg &){      cout <<
        "Hello ";
    });
continue_node< continue_msg > w( g,
    []( const continue_msg &){
        cout << "World\n";
    });
make_edge( h, w );
h.try_put(continue_msg());
g.wait_for_all();
```




# INTEL TBB FLOW GRAPH NODE TYPES

## Functional


source\_node




continue\_node



function\_node



multifunction\_node




## Buffering


buffer\_node




queue\_node



priority\_queue\_node




sequencer\_node




## Split / Join

queueing join




reserving join




tag matching join



split\_node




indexer\_node



## Other

broadcast\_node



write\_once\_node



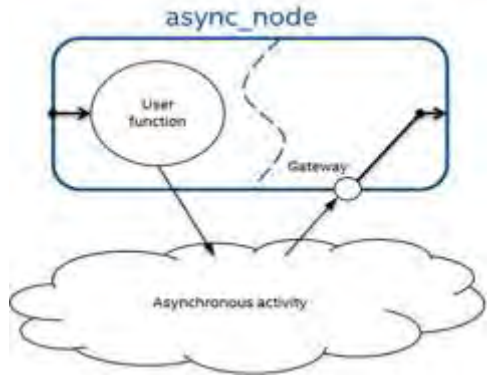
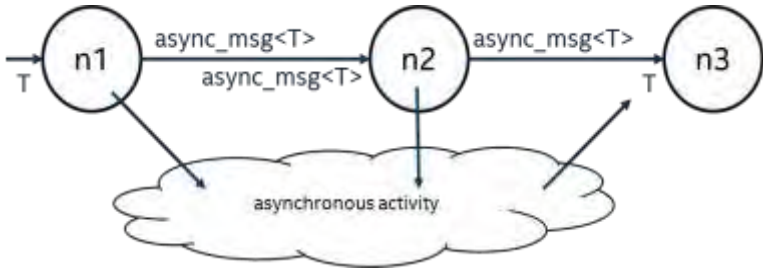
overwrite\_node



limiter\_node



# INTEL TBB FLOW GRAPH FOR HETEROGENEOUS PROGRAMING

Feature	Description	Diagram
<code>async_node&lt;Input,Output&gt;</code> <i>Available as production feature</i>	Basic building block. Enables async communication from a single/isolated node to an async activity. User responsible for managing communication. Graph runs on host.	 <p>The diagram shows a rectangular box labeled 'async_node'. Inside the box, on the left, is a circle labeled 'User function'. On the right side of the box is a small circle labeled 'Gateway'. A dashed line separates the 'User function' from the 'Gateway'. An arrow points from the 'User function' to the 'Gateway'. Below the box is a cloud shape labeled 'Asynchronous activity'. An arrow points from the 'Gateway' to the 'Asynchronous activity'.</p>
<code>async_msg&lt;T&gt;</code> <i>Available as preview feature</i>	Basic building block. Enables async communication with chaining across graph nodes. User responsible for managing communication. Graph runs on the host.	 <p>The diagram shows three circular nodes labeled 'n1', 'n2', and 'n3' arranged horizontally. An arrow labeled 'T' points into 'n1' from the left. An arrow labeled 'T' points out of 'n3' to the right. Between 'n1' and 'n2', and between 'n2' and 'n3', there are two parallel arrows: the top one is labeled 'async_msg&lt;T&gt;' and the bottom one is labeled 'async_msg&lt;T&gt;'. Below the nodes is a cloud shape labeled 'asynchronous activity'. Arrows point from each of the three nodes ('n1', 'n2', and 'n3') down to the 'asynchronous activity' cloud.</p>

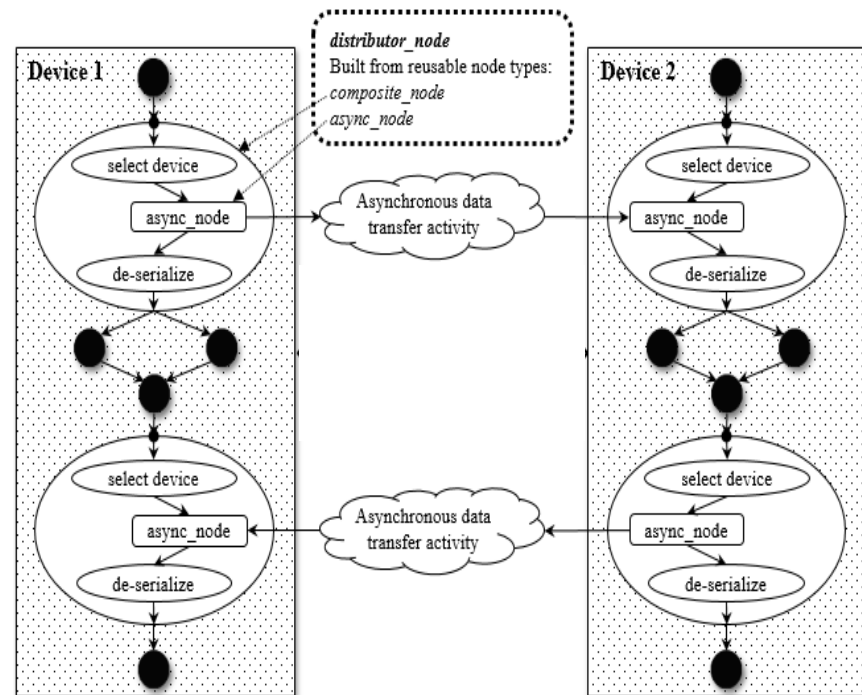
# THE DISTRIBUTOR\_NODE (EXPERIMENTAL)

Enables communication between different memory domains. Each device is capable of running a graph; e.g. hosts, co-processor, etc...

Graphs is execute on all devices


Communication can be initiated from any device to any device

Sub-graphs are executed on a device between communication points



# HETEROGINITY: INTEL FLOW GRAPH AS COORDINATION LAYER

- Exposes parallelism between blocks; simplifies integration
- The glue that connects distributed HW and SW IP
- Libraries implemented using Intel® TBB will compose seamlessly
- Dynamic load balancing
- Supports Distributed memory



# STAC-A2\* - THE FINANCIAL INDUSTRY BENCHMARK


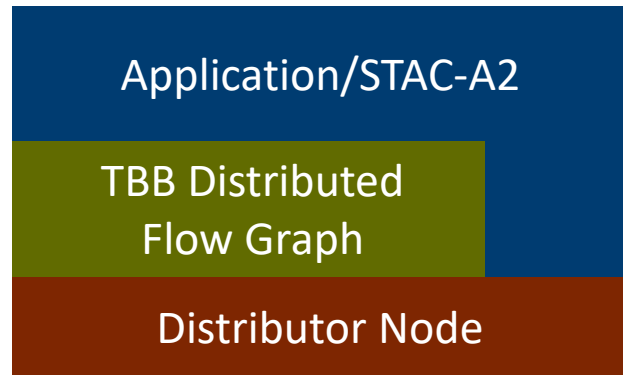



- The STAC-A2 Benchmark is a suite for testing technology platforms (SUT)
- STAC-A2 specifications are based on Market Risk analysis
- “Customers” define the specifications, “Vendors” implement the code
- 7 types of Greeks
- Complex Multi-Asset pricing of American Options
  - Heston Stochastic Volatility underlying price model
  - Least-Squares Monte Carlo / Longstaff-Schwartz algorithm
- Intel first published the benchmark results in Supercomputing'12

\* “STAC” and all STAC names are trademarks or registered trademarks of the Securities Technology Analysis Center LLC.

# STAC-A2\* INTEL IMPLEMENTATION HIGHLIGHTS

- Intel® Parallel Studio components used:
  - Intel Compiler & OpenMP 4.0 for vectorization
  - Intel TBB parallel algorithms for for-join constructs
  - Intel TBB Flow Graph for task distribution
  - Intel® Math Kernel Library (Intel® MKL)
- distributor\_node for distributed compute
- A token-based approach for dynamic load balancing between the main CPU and “coprocessors”



# BREAKING DEPENDENCIES IN MONTE-CARLO (STOCHASTIC VOL)

```
for (unsigned p = 0; p < nPaths; ++p)
{
    double mV[nTimeSteps]; // Spot Volatility state
    double mY[nTimeSteps]; // Spot Price state
    ....
    for (unsigned int t = 0; t < nTimeSteps; ++t){
        double logSpotPrice = func(mY[t], mV[t]); // Backward dependency
        mY[t+1] = logSpotPrice * A[t];
        mV[t+1] = logSpotPrice * B[t] + C[t] * mV[t];
        price[p][t] = exp(mY[t]);
    }
}
```

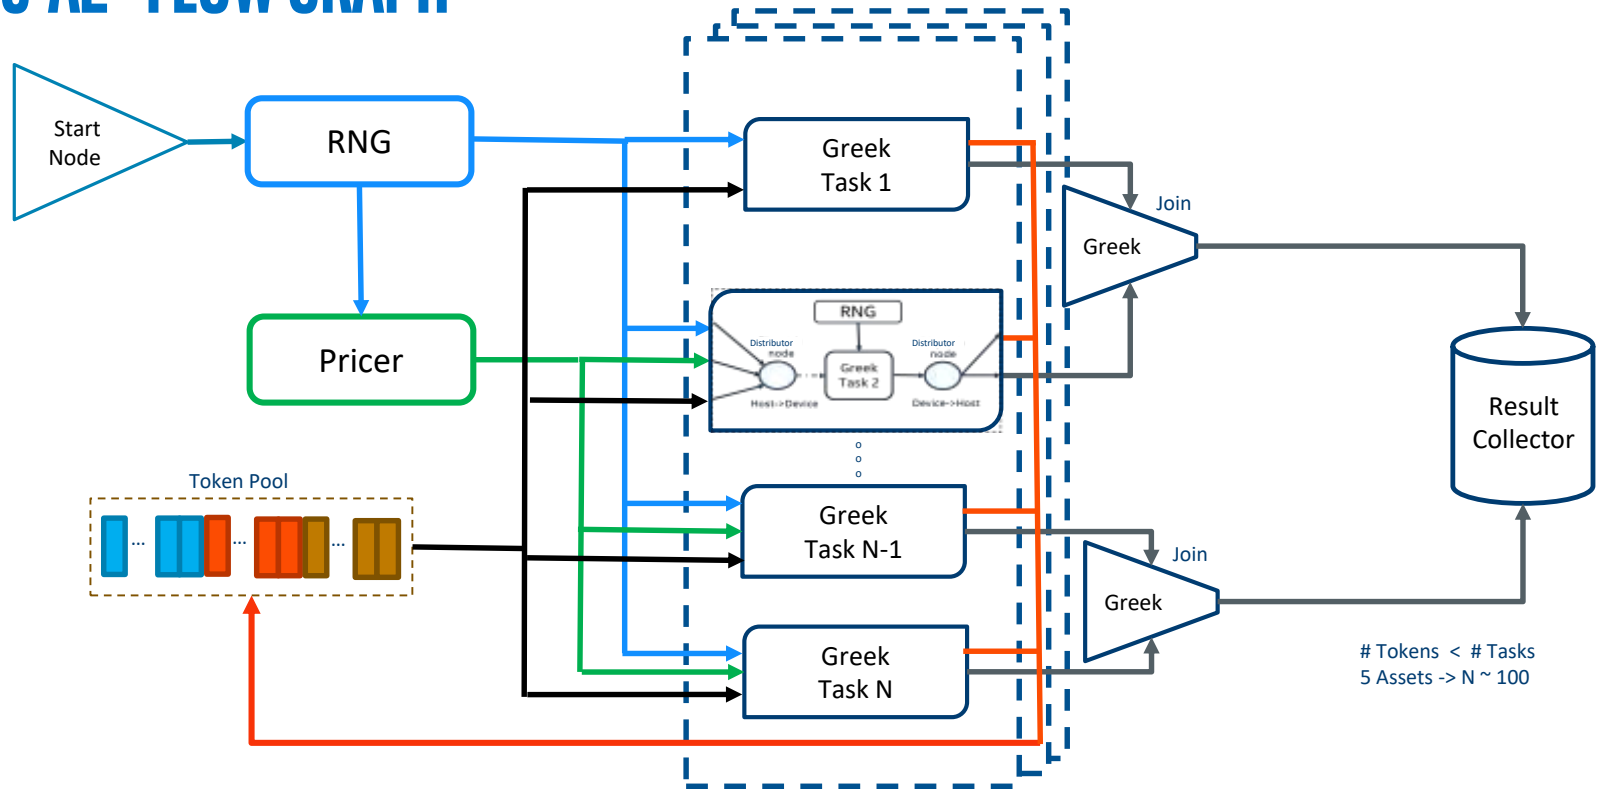


# BREAKING DEPENDENCIES IN MONTE-CARLO (STOCHASTIC VOL)

```
tbb::parallel_for(tbb::blocked_range<int>(0, nPaths, 256),
 [&](const tbb::blocked_range<int>& r) {
    const auto block_size = r.size();
    double mV[nTimeSteps][block_size]; // Spot Volatility state
    double mY[nTimeSteps][block_size]; // Spot Price state
    for (unsigned int t = 0; t < nTimeSteps; ++t){
        for (unsigned p = 0; p < nPaths; ++p)
        {
            double logSpotPrice = func(mY[t][p], mV[t][p]);
            mY[t+1][p] = logSpotPrice * A[t];
            mV[t+1][p] = logSpotPrice * B[t] + C[t] * mV[t][p];
            price[r.begin()+p][t] = exp(mY[t][p]);
        }
    }
}
```

```
for (unsigned p = 0; p < nPaths; ++p)
{
    double mV[nTimeSteps]; // Spot Volatility state
    double mY[nTimeSteps]; // Spot Price state
    .....
    for (unsigned int t = 0; t < nTimeSteps; ++t){
        double logSpotPrice = func(mY[t], mV[t]); // Backward
        dependency
        mY[t+1] = logSpotPrice * A[t];
        mV[t+1] = logSpotPrice * B[t] + C[t] * mV[t];
        price[p][t] = exp(mY[t]);
    }
}
```

# STAC-A2\* FLOW GRAPH




# HETEROGENEOUS GREEK NODE

```
#pragma offload_attribute(push, target(mic))
typedef execution_node < tbb::flow::tuple<std::shared_ptr<GreekResults>, device_token_t>, double> execution_node_theta_t;
...
void CreateGraph(...) {
...
theta_node = std::make_shared<execution_node_theta_t>(_g,
[arena, pWS, randoms](const std::shared_ptr<GreekResults>&, const device_token_t& t) -> double {
    double pv = 0.;
    std::shared_ptr<ArrayContainer<double>> unCorrRandomNumbers;
    randoms->try_get(unCorrRandomNumbers);
    const double deltaT = 1.0 / 100.0;
    pv = f_scenario_adj<false>(pWS->r, ..., pWS->A, unCorrRandomNumbers);
    return pv;
}
, true));
...
}
#pragma offload_attribute(pop)
```

Same code executed on Xeon and Xeon Phi, Enabled by Intel Compiler

# EVOLUTION STAC-A2 - 37X OVER 4 YEARS



	Intel Xeon processor E5 2697-V2	Intel Xeon processor E5 2697-V2	Intel Xeon E5 2697-V2 + Xeon Phi	Intel Xeon E5 2697-V3	Intel Xeon E5 2697-V3+ Xeon Phi	Intel Xeon E5 2697-V3+ 2*Xeon Phi	Intel Xeon Phi 7290	Intel Xeon Phi 7290	Intel Xeon Platinum 8180
	2013	2014	2014	2014	2014	2015	2016	2017	2017
cores	24	24	24+61	36	36+61	36+122	72	72	56
Threads	48	48	48+244	72	72+244	72+488	288	288	112
Vectors(Size)	256	256	256+512	256	256+512	256+2*512	512	512	512
Parallelization	OpenMP	TBB	TBB	TBB	TBB	TBB	TBB	TBB	TBB
Vectorization	#SIMD	OpenMP	OpenMP	OpenMP	OpenMP	OpenMP	OpenMP	OpenMP	OpenMP
Heterogeneity	N/A	N/A	OpenMP	N/A	OpenMP	TBB	N/A	N/A	N/A
Greek time	4.8	1.0	0.63	0.81	0.53	0.216	0.207	0.141	0.128

Appropriate Threading model

1<sup>st</sup> Heterogeneous Implementation

Dynamic Load Balancing between 3 devices

Improved Memory accesses(SDLT)

\*Other names and brands may be claimed as the property of others.  
For more complete information about compiler optimizations, see our [Optimization Notice](#).

U.S. and/or other countries.

Intel Software

# SUMMARY

- Identify 3 levels of parallelism to leverage ALL available HW resources
- Intel TBB allows compasability and utilization of the compute resources
- Intel TBB Flow Graph extensions allow task distribution & dynamic load balancing
- Intel tools support C/C++ for programing in heterogeneous environment

# CODE THAT PERFORMS AND OUTPERFORMS

Download a *free*, 30-day trial of  
Intel® Parallel Studio XE 2018 today

<https://software.intel.com/en-us/intel-parallel-studio-xe/try-buy>

## AND DON'T FORGET...

To check your inbox for the evaluation survey which will be emailed after this presentation.

### P.S.

Everyone who fills out the survey will receive a personalized certificate indicating completion of the training!



# LEGAL DISCLAIMERS


INFORMATION IN THIS DOCUMENT IS PROVIDED “AS IS”. NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as STAC-A2, SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

For more complete information about compiler optimizations, see our Optimization Notice at <https://software.intel.com/en-us/articles/optimization-notice#opt-en>.

© 2017 Intel Corporation. All rights reserved. Intel and the Intel logo are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries.

\*Other names and brands may be claimed as the property of others.



Software